

Lecture 11: Trajectory Optimization

Recap

- Dynamic Programming solve $J^*(x)$
 - on a mesh
 - neural value iteration (less guarantees in terms of convergence)
- LQR ~~(Q,R)~~
 - scales to arbitrarily high dimensions
 - restricted to linear systems
- SOS / Lyapunov methods

Basic Trajectory formulation

$$\dot{x} = f(x, u) \leftarrow \text{dynamical system}$$

Instead of saying our decision variables are somehow like a

Cost-to-go everywhere, we are going to restrict ourselves to just

a trajectory (it is a cont. traj, x is a cont. traj)

for one initial cond.

$$\min_{x(\cdot), u(\cdot)} \int_0^{t_f} l(x(t), u(t)) dt$$

$$\dot{x} = f(x, u), \forall t \text{ (dynamic constraints)}$$

+ additional constraints

$$x(0) = x_0$$

Some class of cont. curves which satisfy all eqns.

Linear discrete-time

Convex Cost func (like linear in the Cost func or a quadratic in the Cost func)

$$\min_{x[\cdot], u[\cdot]} \sum_{n=0}^N l(x[n], u[n])$$

(finite set of decision variables which represent x & u at particular time)

(some finite horizon $n=0 \dots N$)

$$\text{s.t. } x[n+1] = Ax[n] + Bu[n]$$

$$x[0] = x_0$$

+ additional constraints

$$u_{\min} \leq u[n] \leq u_{\max} \forall n$$

input limits are all linear convex constraints

one thing about the linear dynamics is that the dynamics are linear in the optimization problem

Decision variable $x[0], \dots, x[N]$

(Note: $x[0]$ is not a decision variable since it's fixed)

$u[0], \dots, u[N-1]$

"Direct transcription"

So, if we are willing to choose some convex cost function like linear in the cost function or a quadratic in the cost function, then we are in the state where even though we are doing trajectory optimization now, we can still have a convex optimization problem.

Convex costs



$$L(x, u) = x^T Q x + u^T R u, \quad Q \geq 0 \\ R > 0$$

⇒ Quadratic program

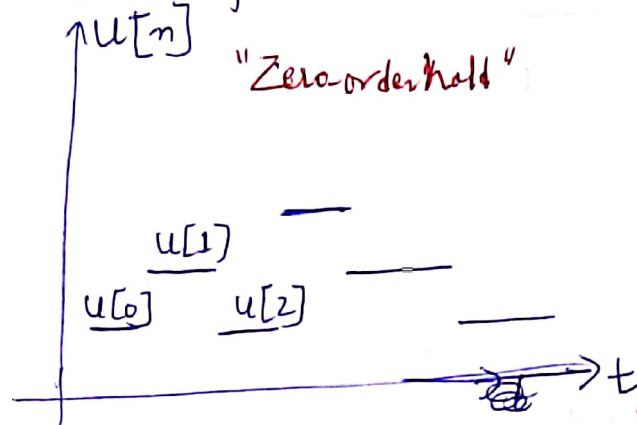
(quadratic cost and linear constraints)
(then we are in the realm of quadratic programming)

Note - If we want to solve the minimum time problem we can actually solve many convex optimization problems, just keep increasing until the first one that succeeds, that's our minimum time.

Discrete time approximation → $x[n+1] = x[n] + \Delta t (A x[n] + B u[n])$
"Euler" Integration



Discrete time optimization



At each time, we are evaluating the dynamics as if u is fixed.

A first-order hold would be linear interpolation.

Linear discrete-time traj. optimization

Looks a lot like LQR but doing something LQR couldn't do. We could put input limits, ~~and~~ state constraints. In LQR, it's hard to solve for input limits and state constraints for all x .

Alternative "Direct Shooting" transcription

Idea - Solve away $x[n]$

$$x[n+1] = Ax[n] + Bu[n]$$

\Downarrow

Given $x[0], u[0], \dots, u[N-1]$

$$x[n] = A^n x[0] + \sum_{k=0}^{n-1} A^{n-k-1} B u[k]$$

$$\min_{u[\cdot]} \sum_{n=0}^N x^T[n] Q x[n] + u^T[n] R u[n]$$

$$\text{s.t. } x[0] = x_0$$

Compare Shooting Vs Direct transcription

- Shooting has less decision variables. — (good).
- Numerics can be worse
 A^n (if n is big)
- Sparsity missing.

Model predictive Control (MPC)

Direct shooting does not give us feedback controller out of the box. If we just solve that once, it only tells us how to go along from one initial condition. But if we start executing the trajectory and there is a disturbance that knocks us off the trajectory then ~~the~~ the initial solution does not tell us what to do.

LQR was better in that sense. It gave us a whole policy. Shooting transcription gave us just a trajectory.

MPC — the idea that if we can solve the optimization reliably and fast enough then just solve it on every time step. Whenever we find ourselves go ahead & solve the optimization problem (like ^{Direct} shooting) and that gives us a feedback controller.

→ MPC solves trajectory optimization on every time step.

→ If we have an optimization trajectory and we want a policy (if we can trust trajectory optimization) then it can give us policy.

Explicit model predictive control

- Explicit model predictive control — That tries to understand what the solution would be for all states, try to turn that into a policy that we pre-compute so that we don't have to solve trajectory optimization online and that field of explicit model predictive control is hugely valuable but more for the theory it gave us than the practice because basically it tells us that the optimal ~~control~~ policy ~~that~~ the cost-to-go gets ~~really~~ complicated really fast.
- Explicit MPC is trying to solve ~~for all initial conditions~~ for the optimal trajectory from all initial conditions for whatever cost functions.


Wouldn't it have been nice if Δt (time step) had been a decision variable in linear discrete-time case (direct transcription)?

$$x[n+1] = x[n] + \overset{\text{time step}}{h} (A x[n] + B u[n])$$

If we make ' h ' a ^{bilinear} decision variable also, then what used to be linear constraint becomes a bilinear constraint and in general, that's not a convex function.

Non-linear Optimization

SNOPT (Sequential Quadratic Programming)



Makes a local quadratic approximation of the function and will use the 2nd order update to go directly to the minimum and then it will take a new quadratic approximation, jump to the minimum.

→ At each step, it takes a quadratic approximation of the cost and a linear approximation of the constraints and it can solve the constrained optimization and then repeat.

Local minima & trajectory optimization

Issue
Sometimes
it gets stuck
in local minimum

Scott Kindersma

trajectory optimization

Atlas

MPC

Trajectory optimization for
dynamic soaring

— Lukas Beyer

Dexterity